



Priced the way each pipeline is actually run, LSD-Flow delivers the same 100-family library for 3.1× fewer reactions — the opposite of what the earlier from-scratch pricing showed, because that regime measured whose chemistry the planner recognised



SCENT sEH vs S3-GFN (500 candidates, 478 routed) · reward > 7.0 · $\tau = 0.5$ · seed 42.
Both arms priced by SPARROW: our readout is the independent SPARROW audit of the same library (our own count-once estimate reads 125, agreeing to 4.8%).
Conservative reading — our 131 vs the baseline's STRONGEST tested configuration (235: MultiAiZ + a diversity-aware selection) = 1.79×.
Asymmetries reported, not matched — pool 500 vs 26,069; surrogate calls ~500 vs 155,764; route planning 2.25 h vs 0.