


Sampling harder does not close the gap

A 15× increase in sampling buys the competitor about a third more molecules — and its enumerated arm outspends us outright while still delivering under half.



All points: reward gate 7.0, $\tau=0.5$, fixed 100-reaction budget (the Logs/062 settings). The competitor's points are a CONTROLLED sweep — one knob (how many molecules it sampled) moves, so they are joined. Ours vary child-policy and pre-filtering together, so they are drawn unjoined: connecting them would assert a curve we did not measure.