


Hub-batching is cheaper per mode for every generator and both targets — but SCENT's headline edge comes from an acquisition only it can run

matrix16 count-once campaign · 30k trajectories / 200 hubs · gates: sEH reward>7.0, DRD2 reward>0.5 · seed 42



Apples-to-apples (all naive): SCENT 1.22x/1.38x vs RxnFlow 1.23x/1.60x. SCENT's free-frag+K20 rows need promoted dynamic-library fragments only it has. Hatched = FragGFN is a CONTROL, not a peer: its "reactions" are fragment attachments and has_route=0. Both FragGFN model versions are shown (Logs/046/050): cap-6 is the corrected model, cap-9 is superseded and kept only because the cap-6 sEH cell is not trained yet, so swapping DRD2 alone would mix versions across blocks. RGFN cells pending.